

CSC411: Assignment 3

Due on Tuesday, March 19, 2017, 11pm

Minh Nguyen (1000468059)

March 23, 2017

Instruction for reproducing the results:

- *naivebayes.py*: codes for part 1,2,3 as well as picking the samples from the data set.
- *logistic.py*: codes for part 4.
- *compare_nb_log.py*: codes for part 6.
- *word2vec.py*: codes for part 7, 8.

The main program for each part is in a separated function `def part[i]()` where `[i]` is the part of the assignment. For example, if you want to run part 8, you should go to file *word2vec.py* and uncomment the line `part8()`, then run the program to see the result.

Please download and extract *review_polarity.tar.gz* into the same folder that you save my codes. This is the raw movie review data set given by the professor.

As part of the submission, I also submitted *report.zip* which consists all of the images used in the report.

Part 1

Dataset descriptions:

The size of the training, validation and test datasets are 800, 200 and 200 respectively. Each set has an equal number of positive and negative reviews. Below is the summary of the number of unique words in each of dataset.

1. Training set: 800 images, 400 postive and 400 negative reviews.
 - Number of unique words in positive reviews: 23303
 - Number of unique words in negative reviews: 21594
 - Total number of unique words: 32689
2. Test set: 200 images, 100 postive and 100 negative reviews.
 - Number of unique words in positive reviews: 10200
 - Number of unique words in negative reviews: 9656
 - Total number of unique words: 15129
3. Validation set: 200 images, 100 postive and 100 negative reviews.
 - Number of unique words in positive reviews: 10855
 - Number of unique words in negative reviews: 9728
 - Total number of unique words: 15687

Example of three key words

We will look at how often the three words “Nice”, “Bad” and “Boring” occur in the training set. “enjoyed”

- Number of occurences of “enjoyed” in the positive reviews: 26
- Number of occurences of “enjoyed” in the negative reviews: 14

“bad”

- Number of occurences of “bad” in the positive reviews: 104
- Number of occurences of “bad” in the negative reviews: 207

“boring”

- Number of occurences of “boring” in the positive reviews: 19
- Number of occurences of “boring” in the negative reviews: 66

It can be seen from the result that the positive word “enjoyed” occurs almost twice times as often in the positive reviews than in negative reviews. In contrast, the negative words such as “boring” and “bad” occur more frequent in negative reviews than in positive reviews. This is because when people writing positive reviews are more likely to use good words like “enjoyed”, “happy”, “fun” to desribe their positive experience with the movie. Similarly, reviewers who had bad experience with the movie are more likely to use the words like “bad” and “boring” in their reviews. Therefore, it is totally feasible to make a prediction whether the review is positive or negative using just some keywords within the review.

Part 2

Tuning parameter m , k :

I ran the Bayesian classifier on the validation set and observe the performance. I started with small values $m = 1$ and $k = 1$, the performance of the classifier was less than 20 % for both negative reviews and positive reviews. This was because multiple terms $P(a_i|class)$ were too small that when adding their logs $\log(P(a_1|class)) + \log(P(a_2|class)) + \dots + \log(P(a_n|class) + \log(0.5))$, the result was a big negative number, so $e^{\log(P(a_1|class)) + \log(P(a_2|class)) + \dots + \log(P(a_n|class) + \log(0.5))} = e^{-\infty} \approx 0$. Because the number was so small that it could not be represented by the system, the OS rounded the number down to 0, this is called arithmetic underflow. Therefore the posterior $P(positivereviews|a_1, a_2, \dots, a_n) = P(negativereviews|a_1, a_2, \dots, a_n) = 0$ made it impossible to do the classification.

To fix the issue, I needed to scale up each of the term $P(a_i|class)$. We have that $P(a_i = 1|class) = \frac{\text{count}(a_i, class) + mk}{\text{count}(class) + k}$, the term k appears in both denominator and numerator but m only shows up in the numerator, so increasing m would only increase the numerator. I kept k constant at 1, and increased m by 50 at a time. As m increased, each term $P(a_i|class)$ is scaled up made it less likely that underflow could happen, increasing the performance of the classifier. When $m \approx 500.0$, the performance on the positive reviews of the validation set was 95%, and the performance on the negative reviews was only 41 %. Any m greater than 500 would make the scaled values to be too big, that is would cause arithmetic overflow. So I stopped increasing m at 500, and started adjusting k .

As k increased, the performance on the negative reviews improved slightly to 51 % from 43 %. However it would also made the performance on the positive reviews decrease to below 90 %. From this point on ward, I just adjusted k and m by small amount until I had the highest average performance between the positive and negative reviews. Below are the recorded performances on the validation, test and train set with $k = 3$ and $m = 460$.

```

-----Validation set performance-----
Performance on positive reviews: 90.0
Performance on negative reviews: 49.0
Overall result: 69.5
5 -----Test set performance-----
Performance on positive reviews: 92.0
Performance on negative reviews: 53.0
Overall result: 72.5
-----Train set performance-----
10 Performance on positive reviews: 93.25
Performance on negative reviews: 65.5
Overall result: 79.375

```

The low performance on the negative reviews is due to the fact that the negative reviews are usually shorter than the positive reviews. This performance can be improved by using a larger training set, the current performance is based on a training set of 800 samples with 32689 unique words.

Part 3

Top 10 words that indicate the review is negative:

if, ?, why, boring, worst, bad, better, it's, script, plot.

Top 10 words that indicate the review is positive:

best, great, also, him, makes, world, while, both, life, young.

Note Professor confirmed on Piazza that there is no need to remove punctuation such as “?”.

Explanation A word is a strong indicator of a positive review if the posterior probability $P(\text{positivereview}|\text{word})$ is high and $P(\text{negativereview}|\text{word})$ is low. On the other hand, if $P(\text{positivereview}|\text{word})$ is low and $P(\text{negativereview}|\text{word})$ is high, then the review with that word is more likely to be negative.

From Bayesian rule, we have that $P(\text{class}|\text{word}) = \frac{P(\text{word}|\text{class})P(\text{class})}{P(\text{word})}$. So $P(\text{positivereview}|\text{word}) \propto P(\text{word}|\text{positivereview})$, and $P(\text{negativereview}|\text{word}) \propto P(\text{word}|\text{negativereview})$ as $P(\text{class}) = 0.5$ is constant since we picked an equal number of samples from each class.

Then if $P(\text{word}|\text{positivereview}) - P(\text{word}|\text{negativereview})$ is high, $P(\text{positivereview}|\text{word}) - P(\text{negativereview}|\text{word})$ is also high. I used Bayesian algorithm to calculate $P(\text{word}|\text{positivereview})$ and $P(\text{word}|\text{negativereview})$ for each word, then I calculated the difference $P(\text{word}|\text{positivereview}) - P(\text{word}|\text{negativereview})$. The top 10 words that have highest $P(\text{word}|\text{positivereview}) - P(\text{word}|\text{negativereview})$ are words that occur frequently in positive reviews but rarely in negative reviews, therefore is a strong positive review's indication. And the bottom 10 words that have least $P(\text{word}|\text{positivereview}) - P(\text{word}|\text{negativereview})$ are strong indications of negative reviews.

Part 4

To choose the right regularization parameter, I observe the performance on the validation set. I first ran the program with $\lambda = 0$, so the high weights would not be penalized at all. The performance on the training set quickly went to 100 % after 30 iterations, whereas the after 300 iterations the performance on the validation set was barely 80 %. Below is the learning curve obtained by running the program with no regularization.




Figure 1: Learning curve with no regularization (i.e. $\lambda = 0$) based on 300 iterations.

The large performance gap between training and test sets indicate that overfitting was occurred. The model was better at learning the specific data in the training set than the general relationship.

To solve the issue, I started setting λ to be some non-zero values. Iteratively, I tried to set λ to be 0.1, 1, 10, etc. and plotted the learning curves. I observed that as λ got higher, the performance on the training set was decreasing while the performance on the validation set was increasing. This closed the performance gap, and therefore reduced the overfitting. At $\lambda = 1000$, the performance on the validation and training sets were 84 % and 94 %, respectively. reduced the number of iterations to 200 from 300 because after 200 iterations, the performance on the validation set was no longer improved. For any $\lambda > 1000$, the performances on all three sets decreased. So $\lambda = 1000$ was the best regularization parameter for my model. Below is the learning curve with the regularization applied.




Figure 2: Learning curve with weight decay ($\lambda = 1000$) based on 200 iterations. Note that the maximum value on the y axis is 95, not 100 like in the previous plot.

Part 5

$$\theta_0 + \theta_1 I_1(x) + \theta_2 I_2(x) + \dots + \theta_k I_k(x) > \text{thr} \quad (1)$$

Naive Bayesian: Let k be the number of unique words in the training set, and x is the document.

$$\begin{aligned} P(\text{class}_{\text{positive}}|a_1, \dots, a_k) &= \frac{P(a_1, a_2, \dots, a_k|\text{class}_{\text{positive}}) \times P(\text{class}_{\text{positive}})}{a_1, a_2, \dots, a_k} \\ \implies P(\text{class}_{\text{positive}}|a_1, \dots, a_k) &\propto P(a_1, a_2, \dots, a_k|\text{class}_{\text{positive}}) \times P(\text{class}_{\text{positive}}) \\ \implies P(\text{class}_{\text{positive}}|a_1, \dots, a_k) &\propto (P(a_1|\text{class}_{\text{positive}}) \times \dots \times P(a_k|\text{class}_{\text{positive}})) \times P(\text{class}_{\text{positive}}) \\ \implies P(\text{class}_{\text{positive}}|a_1, \dots, a_k) &\propto \log(P(a_1|\text{class}_{\text{positive}})) + \dots + \log(P(a_k|\text{class}_{\text{positive}})) + \log(P(\text{class}_{\text{positive}})) \end{aligned} \quad (2)$$

Similarly, we get the following equation for the negative class.

$$P(\text{class}_{\text{negative}}|a_1, \dots, a_k) \propto \log(P(a_1|\text{class}_{\text{negative}})) + \dots + \log(P(a_k|\text{class}_{\text{negative}})) + \log(P(\text{class}_{\text{negative}})) \quad (3)$$

Let take the difference between the two.

$$(2) - (3) = \left(\sum_i^k (\log(P(a_i|\text{class}_{\text{positive}})) - \log(P(a_i|\text{class}_{\text{negative}}))) \right) + (\log(P(\text{class}_{\text{positive}})) - \log(P(\text{class}_{\text{negative}}))) \quad (4)$$

If $(2) - (3) > 0$, it means that if **all** the unique words in the training set appear in the document, then it is more likely that the document is a positive review. Otherwise, if $(2) - (3) < 0$, there is higher chance that the document is a negative review. We can think of each term $(\log(P(a_i|\text{class}_{\text{positive}})) - \log(P(a_i|\text{class}_{\text{negative}})))$ in the summation as θ_i . However, **not all** key words a_i will be showed up in the test document. So we need an indicator function $I_i(x)$ that indicates if the key word a_i is in the document x . If a_i is in x , then $I_i(x) = 1$, otherwise $I_i(x) = 0$. Adding in the indicator function into the equation (4), we have the following equation for doing classification:

$$\begin{aligned} &= \left(\sum_i^k ((\log(P(a_i|\text{class}_{\text{positive}})) - \log(P(a_i|\text{class}_{\text{negative}}))) I_i(x)) \right) + (\log(P(\text{class}_{\text{positive}})) - \log(P(\text{class}_{\text{negative}}))) \\ &= \left(\sum_i^k (\theta_i \times I_i(x)) \right) + (\log(P(\text{class}_{\text{positive}})) - \log(P(\text{class}_{\text{negative}}))) \\ &= \left(\sum_i^k (\theta_i \times I_i(x)) \right) + \theta_0 \end{aligned} \quad (5)$$

Logistic Regression Assuming that we use one-hot encoding, we get two sets of weights, one for each of the two classes. This is similar to the two set of probabilities we got in Naive Bayesian algorithm.

$$\begin{aligned} W_{\text{positivereview}} &= \{W_{\text{word1},+}, W_{\text{word2},+}, \dots, W_{\text{wordk},+}\} \\ W_{\text{negativereview}} &= \{W_{\text{word1},-}, W_{\text{word2},-}, \dots, W_{\text{wordk},-}\} \end{aligned}$$

We can define the classification function as follow:

$$\begin{aligned} &= \left(\sum_i^k ((W_{\text{wordi},+} - W_{\text{wordi},-}) I_i(x)) \right) + (W_{\text{bias}+} - W_{\text{bias}-}) \\ &= \left(\sum_i^k (\theta_i \times I_i(x)) \right) + \theta_0 \end{aligned} \quad (6)$$

Similar to what we got with Bayesian classification, $I_i(x)$ is the indicator function that outputs 0 or 1 to indicate whether the key word i is in the document x . θ_i is $(W_{wordi,+} - W_{wordi,-})$ and θ_0 is the weight bias.

Note: $+$ denotes the class positive review, and $-$ denotes the class negative review.

Part 6

Top 50 θ for positive reviews with Naive Bayesian classification:

```
{'own': 0.030833333333333268, 'family': 0.02416666666666667, 'simple': 0.0266666
6666666684, 'years': 0.03750000000000009, 'see': 0.02416666666666667, 'human': 0
.03416666666666668, 'seen': 0.03166666666666673, 'subtle': 0.03249999999999997,
'still': 0.026666666666666616, 'yet': 0.03499999999999992, 'best': 0.05416666666
6666696, 'perfect': 0.02499999999999991, 'different': 0.02416666666666667, 'thro
ughout': 0.026666666666666616, 'takes': 0.03166666666666651, 'true': 0.030833333
333333268, 'young': 0.03833333333333333, 'also': 0.05416666666666696, 'does': 0
.02416666666666667, 'between': 0.03583333333333384, 'performance': 0.0266666666
66666616, 'several': 0.03166666666666673, 'life': 0.0666666666666665, 'attentio
n': 0.02416666666666667, 'may': 0.03499999999999992, 'each': 0.03166666666666673
, 'most': 0.03749999999999987, 'job': 0.02416666666666667, 'others': 0.032499999
99999997, 'world': 0.05499999999999994, 'strong': 0.029166666666666563, 'friends
': 0.03166666666666673, 'him': 0.05083333333333351, 'begins': 0.0283333333333333
2, 'both': 0.04333333333333346, 'great': 0.044166666666666465, 'perfectly': 0.0
27499999999999858, 'especially': 0.033333333333333215, 'effective': 0.0258333333
33333375, 'always': 0.030833333333333268, 'brilliant': 0.030833333333333268, 'ma
ny': 0.03749999999999987, 'allows': 0.02416666666666667, 'will': 0.0283333333333
3332, 'while': 0.03916666666666657, 'performances': 0.03500000000000014, 'very':
0.03333333333333344, 'american': 0.030000000000000027, 'makes': 0.0425000000000
00204, 'almost': 0.0250000000000000133}
```

Top 50 θ for negative reviews with Naive Bayesian classification:

```
{'just': -0.025833333333333153, 'give': -0.023333333333333206, "didn't": -0.0233
33333333333428, 'anyway': -0.0225000000000000187, 'wasted': -0.026666666666666616
, 'awful': -0.0250000000000000133, 'looks': -0.030833333333333268, 'if': -0.04750
000000000001, '!': -0.03249999999999997, 'plot': -0.03833333333333333, 'guess': -
0.026666666666666616, 'i': -0.026666666666666684, 'script': -0.04833333333333334,
"it's": -0.0400000000000000036, 'least': -0.03499999999999992, 'either': -0.0233
33333333333206, 'better': -0.03833333333333333, "wasn't": -0.0250000000000000133,
"there's": -0.02833333333333332, 'got': -0.023333333333333206, 'waste': -0.0250
000000000000133, 'decent': -0.022499999999999964, '?': -0.06916666666666666, 'supp
osed': -0.036666666666666625, 'poor': -0.02416666666666667, 'worse': -0.02416666
66666667, 'mess': -0.026666666666666616, 'material': -0.02499999999999991, 'rea
son': -0.02833333333333332, 'ridiculous': -0.02416666666666667, 'worst': -0.0549
9999999999994, 'nothing': -0.027499999999999858, 'why': -0.04083333333333328, 'l
ook': -0.025833333333333375, 'none': -0.027499999999999858, 'save': -0.022499999
999999964, 'talent': -0.0225000000000000187, 'unfortunately': -0.03333333333333334
4, 'course': -0.022499999999999964, 'did': -0.027500000000000008, 'maybe': -0.033
33333333333215, 'lame': -0.02499999999999991, 'terrible': -0.023333333333333428
, 'thing': -0.029999999999999805, 'bad': -0.08583333333333343, 'stupid': -0.0341
6666666666668, "can't": -0.03166666666666673, 'boring': -0.039166666666666794, '
predictable': -0.02833333333333332, "isn't": -0.02833333333333332}
```

Top 50 θ for positive reviews with Logistic Regression classification:

```
{'own': 0.013312923, 'simple': 0.012355411, 'years': 0.016862687, 'see': 0.01202
184, 'human': 0.015253512, 'subtle': 0.015125709, 'seen': 0.016072426, 'still':
0.012846986, 'yet': 0.016108947, 'best': 0.025253143, 'perfect': 0.011767382, 't
hroughout': 0.012045614, 'takes': 0.013620302, 'true': 0.01362077, 'young': 0.01
529683, 'terrific': 0.011474546, 'also': 0.025604311, 'between': 0.015466765, 'p
```

```

performance': 0.011917949, 'several': 0.014882077, 'life': 0.029041685, 'attentio
n': 0.011084631, 'may': 0.014872752, 'very': 0.016251992, 'each': 0.013253912, '
overall': 0.012112376, 'hilarious': 0.011038193, 'most': 0.016242225, 'job': 0.0
11477605, 'others': 0.015621452, 'world': 0.024853006, 'strong': 0.012559932, 'f
10 riends': 0.016162531, 'him': 0.023612441, 'begins': 0.012439251, 'both': 0.01898
8151, 'great': 0.020791847, 'perfectly': 0.013204739, 'especially': 0.016167773,
'effective': 0.011887925, 'always': 0.013800211, 'brilliant': 0.015066012, 'per
formances': 0.015978094, 'allows': 0.012258044, 'will': 0.012324894, 'while': 0.
017829601, 'many': 0.015284918, 'american': 0.013127217, 'makes': 0.020303458, '
15 almost': 0.011035507}

```

Top 50 θ for negative reviews with Logistic Regression classification:

```

{"isn't": -0.012588516, 'just': -0.01103092, 'give': -0.011644898, 'wasted': -0.
013344445, 'awful': -0.012129982, 'looks': -0.015486099, 'have': -0.010863103, '
if': -0.024498831, '!': -0.013051623, 'plot': -0.018612798, 'guess': -0.01262189
4, 'script': -0.024369545, 'it's': -0.018512912, 'least': -0.016362293, 'better'
5 : -0.019479303, 'wasn't': -0.012312988, 'there's': -0.012758968, 'waste': -0.012
352165, '?': -0.032902829, 'either': -0.011666148, 'poor': -0.012486506, 'then':
-0.010981265, 'worse': -0.01204066, 'mess': -0.013252348, 'material': -0.013801
686, 'director': -0.012016099, 'reason': -0.014537231, 'ridiculous': -0.01187721
10 5, 'worst': -0.027064471, 'nothing': -0.01420748, 'attempt': -0.011391882, 'why'
: -0.019721836, 'look': -0.01285463, 'none': -0.013958635, 'save': -0.01141117,
'talent': -0.011549305, 'unfortunately': -0.016719311, 'did': -0.012901056, 'may
be': -0.016109418, 'could': -0.010828576, 'lame': -0.01203445, 'terrible': -0.01
15 1574896, 'thing': -0.013621684, 'bad': -0.040762775, 'stupid': -0.015774131, 'su
pposed': -0.018499222, 'boring': -0.018978249, 'predictable': -0.014013186, 'or'
: -0.010978724, 'can't': -0.014851704}

```

Explanation:

I used the definition of θ in part 5 to do this question. The 50 highest θ s are corresponding to words that better predict positive reviews, and the 50 lowest θ s are corresponding words that are useful to predict negative reviews. Let $S_{\text{bayesian},+}$, $S_{\text{logistic},+}$, $S_{\text{bayesian},-}$, $S_{\text{logistic},-}$ be the sets of words corresponding to the 4 sets of θ s obtained for positive and negative reviews using the two classifiers.

$S_{\text{bayesian},+}$ and $S_{\text{logistic},+}$ have 47 out of 50 words in common. Positive words like *great*, *perfect*, *effective*, *best*, *brilliant*, etc. appear in both sets.

$S_{\text{bayesian},-}$ and $S_{\text{logistic},-}$ share 45 out of 50 words in each set. Common negative words for commenting such as *worst*, *waste*, *ridiculous*, *bad*, *stupid*, etc. show up in both sets.

These similarities suggest that the performance of the two classifiers on the testing set should be similar. Moreover, it also verifies that our implementations of the two classifiers are correct.

Part 7

Experiment Description:

I used the words in the original reviews to build the training and testing datasets. My program randomly selected a set of 30 documents (15 positive reviews and 15 negative reviews) and another set of two documents (1 positive review and 1 negative review) in which the words are used to build the training and testing sets, respectively. In each document, the program collected every pair of words that are next to each other. For every adjacent words pair, the program also **randomly** picked a non-adjacent words pair in the document and added its embedding vectors to the dataset along with the embedding vectors of the adjacent words pair. Therefore, the number of adjacent and non-adjacent words pairs in each dataset was equal. Below is the pseudo code which explains how the word pairs are picked, so that we always have an equal number of samples in each class.

```
documents = randomly pick set of documents to get the words
for each doc in documents:
    list_of_words = collect words in the doc.
    for word in list_of_words:
        ignore words that are do not have vector mapping.
        add (word, word_behind) into the dataset if word_behind is available.
        add (word, word_infront) into the dataset if word_infront is available.

    also add (word, random word in doc) into the dataset for each
        (word, word_behind)/(word, word_infront) added.
```

The two classes for classification were *adjacent* and *non-adjacent*. Using one-hot encoding, the adjacent words pair was labelled [1, 0], and the non-adjacent words pairs was labelled [0, 1]. The input to the logistic regression model was the concatenation of the two 128-bits embedding vectors of two words in a pair. So in total, there were 256 bits input for each pair of adjacent/non-adjacent words.

Learning Curve:

Below is the learning curve obtained from running logistic regression for 100 iterations:




Figure 3: Learning curve

With 67420 pairs of words in the training set, and 7108 pairs of words in the test set, the classifier correctly classified about 80 % of the samples in both training and testing sets.

Part 8

I used Ecludian distance to calculate the difference between the vectors. The two additional examples that I have are for the words *best* and *very*.

10 words that are related to the word *good* based on the embeddings are:

perplexing reinforcing manipulate wonderful admiral great bad funny decent underused

10 words that are related to the word *story* based on the embeddings are:

ricci film acclaim plot domineering benito simmer lift interviews sitter

10 words that are related to the word *very* based on the embeddings are:

extremely fairly somewhat pretty skipping too sailboat quite oft bit

10 words that are related to the word *best* based on the embeddings are:

top kyle pounder simplicity popular origins worst neal morton grahams